

# AI full-stack intern

Synthetic fixture · no real personal data

|                                                                                                                                                                                                                                                                                                                                                                                                                        |                                          |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------|
| Name / Education / Skills<br>AI full-stack intern<br>Source: synthetic benchmark fixture; no real personal data.<br>Role: AI full-stack intern<br>Requirement 1: Deliver Vue and Java features end to end.<br>Requirement 2: Stream model output with SSE and preserve session state.<br>Requirement 3: Keep API p95 below 800 ms for non-model endpoints.<br>Constraint: p95 = the 95th percentile of request latency | Column B<br>Experience and project notes |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------|

Source: synthetic benchmark fixture; no real personal data.

Role: AI full-stack intern

Requirement 1: Deliver Vue and Java features end to end.

Requirement 2: Stream model output with SSE and preserve session state.

Requirement 3: Keep API p95 below 800 ms for non-model endpoints.

Constraint: p95 = the 95th percentile of request latency